

Integrating an Econometric Demand Model with Deep Q-Network Pricing: A Simulation Study

[Author Name(s)]

[Department / School]

[University Name]

City, Country

email@institution.edu

Abstract—Retail dynamic pricing requires both demand structure and sequential adaptivity; purely econometric and purely model-free reinforcement-learning (RL) methods address complementary aspects of this problem. We describe *EG-DQN* (Econometric-Guided Deep Q-Network), a hybrid agent that incorporates Generalized Linear Model (GLM) demand signals into a Deep Q-Network through state augmentation, econometric reward shaping, and early action masking. We evaluate the design in a controlled simulation of three substitute products over 1,460 days of synthetic daily demand, with a 120-day holdout on Item 1 and comparison of fixed-price, rule-based, GLM-optimal, pure DQN, and EG-DQN policies. On the holdout window, GLM-optimal pricing and EG-DQN attain identical cumulative revenue and select the same price on the discrete grid; a one-sided Mann-Whitney test does not indicate higher daily revenue for EG-DQN relative to the GLM benchmark ($p = 0.50$). Relative to fixed-price and rule-based baselines, EG-DQN yields higher holdout revenue in this simulation (Mann-Whitney $p < 0.001$ in both comparisons). Pure DQN attains lower holdout revenue than GLM-optimal and EG-DQN under the learned policies. An ablation suggests that econometric reward shaping accounts for the principal holdout improvement over pure DQN, while the full hybrid does not exceed a reward-only configuration on test revenue. Training-episode rewards are modestly higher for EG-DQN than for pure DQN. These findings are confined to a synthetic, per-item, discrete-action environment with partial alignment between the data-generating process and the GLM specification; operational deployment is not claimed.

Index Terms—Dynamic pricing, reinforcement learning, deep Q-network, generalized linear model, price elasticity, demand forecasting, simulation study, retail pricing.

I. INTRODUCTION

Retail managers set prices under incomplete information. When tomorrow’s price is chosen, competitor reactions, seasonal shifts, and the price sensitivity of the current customer pool are not directly observable. Mistakes are costly: a price set too low forfeits revenue that cannot be recovered, and a price set too high suppresses demand immediately.

The literature has developed two complementary approaches. Econometric methods estimate price-response relationships from historical transactions, yielding interpretable parameters—in particular, own-price elasticity—that managers can audit and reconcile with demand theory [1]. Their optimisation is typically open-loop and rests on assumptions such as relatively stable demand relationships over the estimation

window. Reinforcement learning (RL), by contrast, treats pricing as a sequential decision problem in which an agent learns a policy from repeated market interaction [3]. RL adapts to feedback without committing to a fixed parametric form, yet retail applications must contend with sample-efficient learning, noisy rewards, and exploration over large price sets [11], [13].

Apte et al. [11] show that tabular Q-learning can improve on static pricing in simulated retail settings. Groeneveld et al. [3] report that RL remains competitive under demand-model misspecification but can learn more slowly when demand structure is known—the regime in which econometric guidance may be most informative. When live experimentation is impractical, simulation-based evaluation offers a controlled alternative [15], [16]. We adopt that evaluation mode here; operational deployment is not claimed.

We describe *EG-DQN* (Econometric-Guided Deep Q-Network), which couples a Generalized Linear Model (GLM) demand specification with a Deep Q-Network (DQN) pricing agent. GLM-derived signals enter the agent through three mechanisms: state augmentation with the GLM-optimal price at each timestep, econometric reward shaping, and early action masking during exploration. The design is assessed in a synthetic environment with three substitute products, 1,460 days of daily demand, and a 120-day holdout on Item 1; we compare fixed-price, rule-based, GLM-optimal, pure DQN, and EG-DQN policies using holdout revenue, Mann-Whitney tests, and a three-configuration ablation.

Contributions

- We **describe** EG-DQN and three mechanisms for injecting GLM demand signals—optimal-price guidance, reward shaping, and masked exploration—into DQN-based retail pricing.
- We **implement** a reproducible simulation benchmark with known elasticities, five baseline strategies, holdout evaluation on Item 1, non-parametric significance testing, and ablation of econometric reward shaping.
- We **report** that GLM-optimal and EG-DQN policies coincide on holdout revenue in this simulation while exceeding naive fixed-price and rule-based strategies; econometric reward shaping aligns RL outcomes with

GLM-guided prices, and EG-DQN attains modest training-phase gains over pure DQN, without claiming holdout dominance over econometric pricing alone.

II. RELATED WORK

A. Dynamic Pricing and RL

Chenavaz and Dimitrov [2] review how artificial intelligence methods have entered dynamic pricing research across operations and applied economics. Kopalle et al. [1] summarise managerial uses of dynamic pricing and the econometric questions that remain open in retail practice. On the algorithmic side, Groeneveld et al. [3] compare reinforcement learning with data-driven dynamic programming in finite-horizon pricing markets. Liu et al. [4] extend deep RL to joint omni-channel pricing and inventory control when demand is uncertain. Li and Zheng [5] study pricing under inventory constraints and external information within an online learning framework—a reminder that demand estimation and sequential price setting are seldom separable in applied settings.

B. Demand Forecasting for Pricing

Pricing rules draw on demand forecasts, elasticity estimates, or both. Zhang and Meng [6] model time-varying price elasticity from retail time series, relaxing the constant-elasticity assumption embedded in many static GLM specifications. Bu et al. [7] analyse offline demand learning when sales are censored by stockouts, a common complication in retail transaction data. Xia et al. [8] generate synthetic retail demand with known elasticities so that estimated parameters can be checked against ground truth; we adopt that validation logic in our simulation benchmark.

C. Econometric Demand Models

For count demand, a log-linked GLM remains a standard specification when managers require interpretable own- and cross-price effects [6]. Safonov [10] contrasts neural demand estimators with econometric models in retail and notes that flexible machine-learning fit can underuse sparse within-product price variation that a GLM isolates. The implication for hybrid pricing is to treat econometric structure as an auditable input to learning, not as the only admissible policy.

D. RL for Pricing

Apte et al. [11] apply tabular Q-learning to retail pricing without an explicit structural demand block in the agent. Alamdar and Seifi [12] extend deep Q-learning to joint ordering and pricing decisions. Cheung and Simchi-Levi [13] study nonstationary RL and the cost of unrestricted early exploration—one motivation for restricting actions while a demand model is still imprecise. Zheng et al. [14] combine learned demand representations with a dual-agent deep RL policy for pricing and replenishment; their demand-learning layer is the closest architectural neighbour to our GLM-guided agent. When live price experiments are costly or infeasible, offline and simulation-based evaluation protocols offer a practical alternative [15].

Prior retail pricing studies typically pursue either econometric optimisation from a fixed demand model or end-to-end RL with limited structural transparency. Few papers specify separate channels—state information, reward shaping, and action restrictions—through which econometric estimates enter an RL policy, then compare the resulting hybrid against GLM-only and pure DQN baselines on a holdout window with known elasticities. EG-DQN is intended to document that integration in simulation; we do not present it as a deployment-ready pricing system.

III. PROBLEM FORMULATION

A. Market Structure

We consider a single retailer managing $n = 3$ substitute products over a discrete horizon of T days, following multi-product retail pricing formulations [17]. At each period $t \in \{1, \dots, T\}$, the retailer observes state information and sets a price for the focal product from a finite action space. Cross-product substitution is captured in the demand system below; RL policies are estimated *per product*, with rival prices entering as inputs to the econometric model rather than as joint actions in a single multi-product MDP. The objective is a pricing policy $\pi : \mathcal{S} \rightarrow \mathcal{A}$ that maximises expected cumulative revenue:

$$\pi^* = \arg \max_{\pi} \mathbb{E} \left[\sum_{t=1}^T p_t \cdot D_t(p_t, s_t) \right], \quad (1)$$

where p_t and D_t denote the price and realised demand of the focal item on day t .

B. Demand Data-Generating Process

Simulated demand follows a log-linear multiplicative structure grounded in demand theory [1], [6]:

$$D_{it} = B_i(t) \cdot S_i(t) \cdot \prod_{j=1}^n \left(\frac{p_{jt}}{\bar{p}_j} \right)^{\varepsilon_{ij}} + \eta_{it}, \quad (2)$$

where $B_i(t)$ is a time-varying base-demand component (trend), $S_i(t)$ is a multiplicative seasonal index, ε_{ij} is the cross-price elasticity of product i with respect to product j 's price (own-price elasticity when $i = j$), $\bar{p}_j$ is a reference price, and $\eta_{it} \sim \mathcal{N}(0, \sigma_i^2)$ is i.i.d. Gaussian noise. Because $\log \mathbb{E}[D_{it}]$ is linear in $\log p_{jt}$, the Poisson GLM with log-link used later is *approximately* aligned with the simulator; it is not treated as a correctly specified model for the noise term or for all trend-seasonal components.

C. Pricing as a Markov Decision Process

The focal product's pricing problem is written as an MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$.

State space $\mathcal{S}$: $s_t \in \mathbb{R}^5$ comprises: (i) previous-period price; (ii) previous-period observed demand; (iii) day-of-year sine encoding annual seasonality; (iv) normalised time index t/T ; and (v) the GLM-optimal price $p_{\text{GLM},t}^*$. Component (v) is the econometric signal that distinguishes EG-DQN from pure DQN.

Action space $\mathcal{A}$: A discrete grid of $K = 20$ prices on $[p_{\min}, p_{\max}]$. During early training, EG-DQN restricts feasible actions to a neighbourhood of $p_{\text{GLM},t}^*$ (action masking).

Reward — Pure DQN:

$$r_t = p_t \cdot \max(0, \hat{D}_t(p_t, s_t)), \quad (3)$$

where $\hat{D}_t$ is the GLM demand forecast used as a revenue proxy.

Reward — EG-DQN:

$$r_t^{\text{EG}} = r_t + \lambda \cdot \exp\left(-\frac{|p_t - p_{\text{GLM},t}^*|}{\sigma_p}\right), \quad (4)$$

with econometric weight λ and price scale σ_p . The bonus attracts the agent toward GLM-endorsed prices without fixing the selected price.

Discount factor: $\gamma = 0.99$, consistent with daily pricing over finite training episodes.

IV. ECONOMETRIC DEMAND MODEL

A. Specification

We estimate a separate Poisson GLM with log-link for each product i . Count demand is natural for daily units sold; the log-link yields interpretable own- and cross-price elasticities of the form common in retail pricing research [1], [6]. For product i ,

$$\begin{aligned} \log(\mathbb{E}[D_{it}]) = & \alpha_0 + \underbrace{\alpha_1 \log p_{it}}_{\text{own elasticity}} + \sum_{j \neq i} \beta_j \log p_{jt} \\ & + \gamma_1 \text{trend}_t + \gamma_2 \sin\left(\frac{2\pi \text{day}_t}{365}\right) + \gamma_3 \text{weekend}_t, \end{aligned} \quad (5)$$

where α_1 is the own-price elasticity, β_j are cross-price elasticities, and γ_k capture trend, annual seasonality, and weekend effects. Coefficients are estimated by maximum likelihood on the training window only; HC3-robust standard errors are reported. Holdout forecasts and GLM-optimal prices use this fixed fit—we do not refit online on the 120-day test window. As noted in Section III, the log-linear DGP is only approximately aligned with (5); time-varying elasticity is outside the present specification [6].

B. Derived Econometric Signals

Three GLM outputs feed the hybrid agent [10]:

- 1) **Estimated elasticity** $\hat{\varepsilon}_i = \hat{\alpha}_1$. In simulation we compare $\hat{\varepsilon}_i$ to the known DGP own-price elasticity as an *elasticity recovery* check, not as a claim of structural identification.
- 2) **GLM-optimal price** $p_{\text{GLM},t}^*$. For each day t and covariate vector $\mathbf{x}_t$, we evaluate expected revenue $R(p) = p \cdot \hat{D}_i(p, \mathbf{x}_t)$ on the discrete price grid and set $p_{\text{GLM},t}^* = \arg \max_p R(p)$. This quantity enters the EG-DQN state, the early action mask, and the econometric reward bonus.
- 3) **Forecast demand** $\hat{D}_{it}$. The fitted mean demand at the chosen price serves as a revenue proxy in (3), limiting the effect of day-to-day demand noise on the learning signal [7].

Estimation accuracy and holdout forecast metrics are deferred to the experimental results; this section defines only the signals passed to the DQN.

V. REINFORCEMENT LEARNING FRAMEWORK

A. Network Architecture

Each product is assigned its own Deep Q-Network agent. The policy network $Q(s, a; \theta)$ maps the five-dimensional state vector defined in Section III to a Q-value for each of $K = 20$ discrete candidate prices. The network has two hidden layers of 128 units with ReLU activations and Glorot initialisation; parameters are updated in NumPy without a deep-learning framework dependency, which keeps the simulation reproducible on standard hardware. Retail Q-learning studies motivate discrete-action pricing agents of this form [11], [12]. Temporal-difference targets follow the standard DQN backup with a separate target network θ^- ; the regression loss is Huber with $\delta = 1$ rather than squared error, which limits the influence of occasional high-revenue transitions:

$$\mathcal{L}(\theta) = \mathbb{E}_{\mathcal{B}} \left[\ell_{\delta} \left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right) \right]. \quad (6)$$

B. Training Procedure

Training proceeds for 300 episodes per product. At each environment step the agent stores transitions (s_t, a_t, r_t, s_{t+1}) in a replay buffer of capacity 10,000; a minibatch of 64 transitions is drawn uniformly for each gradient update. The target network weights are copied from the online network every 15 episodes. Action selection is ε -greedy with ε decaying from 1.0 to 0.05 over training; gradient norms are clipped at $\ell_2 = 1.0$. The learning rate is 10^{-3} . Pure DQN uses the reward in (3); EG-DQN variants and econometric hyperparameters (λ , masking horizon) are specified in the next subsection.

C. EG-DQN Integration Mechanisms

EG-DQN extends the base DQN with three channels through which the fixed GLM fit from Section IV enters learning. Pure DQN omits all three; the ablation in the results section toggles reward shaping while holding the remaining protocol fixed.

- **State augmentation.** The state vector includes $p_{\text{GLM},t}^*$ at every step (component (v) in Section III), giving the agent an explicit econometric price benchmark alongside lagged market variables.
- **Reward shaping.** Training rewards follow (4) with $\lambda = 30.0$ and σ_p set to the standard deviation of the discrete price grid. The exponential bonus is largest when the selected price coincides with $p_{\text{GLM},t}^*$ and decays smoothly with price distance.
- **Early action masking.** For episodes $e \leq 30$, feasible actions are restricted to grid points within $\pm \sigma_p$ of $p_{\text{GLM},t}^*$, limiting exploratory prices while the replay buffer is still sparse. The full grid is available from episode 31 onward. Restricting early exploration in nonstationary settings can reduce the cost of uninformed trial prices [13].

Algorithm 1 summarises one training episode; the outer loop runs $E = 300$ episodes per product with the hyperparameters in Section V.

Algorithm 1 EG-DQN Training (single episode)

Require: Fitted GLM $\mathcal{G}$, environment $\mathcal{M}$, replay buffer $\mathcal{B}$, weights θ , target θ^- , exploration rate ε

- 1: $s_0 \leftarrow \text{RESET}()$; compute $p_{\text{GLM},0}^*$ from $\mathcal{G}$
- 2: **for** step $t = 0, 1, \dots$ until episode end **do**
- 3: **if** current episode $e \leq 30$ **then** $\mathcal{A}_t \leftarrow \{p \in \text{grid} : |p - p_{\text{GLM},t}^*| \leq \sigma_p\}$
- 4: **else** $\mathcal{A}_t \leftarrow$ full price grid
- 5: **end if**
- 6: $a_t \leftarrow \varepsilon\text{-greedy}(Q(s_t, \cdot; \theta), \mathcal{A}_t)$
- 7: Execute a_t ; observe r_t^{EG} from (4)
- 8: $s_{t+1} \leftarrow$ next state; store $(s_t, a_t, r_t^{\text{EG}}, s_{t+1})$ in $\mathcal{B}$
- 9: **if** $|\mathcal{B}| \geq 64$ **then**
- 10: Sample minibatch from $\mathcal{B}$; update θ via (6) with gradient clip
- 11: **end if**
- 12: $s_t \leftarrow s_{t+1}$
- 13: **end for**
- 14: Decay ε ; if $e \bmod 15 = 0$ then $\theta^- \leftarrow \theta$

VI. SIMULATION ENVIRONMENT

A. Design and Calibration

We evaluate EG-DQN in a controlled synthetic market rather than a live retail test. Simulation allows known ground-truth elasticities, a fixed holdout protocol, and reproducible comparison of pricing policies under identical demand trajectories—an evaluation mode used for offline and retail RL studies when field experimentation is costly or infeasible [8], [9], [15], [16]. We do not treat the synthetic setting as a substitute for real-store validation; it is a design and benchmark study.

The simulator generates daily demand for $n = 3$ substitute products over $T = 1,460$ days (1 January 2021 – 30 December 2024). The final 120 days form an out-of-sample holdout; the main revenue comparison in the results section uses Item 1 on that window. Demand parameters are calibrated to statistical properties typical of large retail benchmarks (M5 / Walmart-style series and related synthetic generators [8]): prices on a \$6–\$18 range, demand coefficient of variation of approximately 0.25, and clear annual seasonality. All stochastic draws and training runs use seed 42.

Environment components are as follows.

- **Trend.** Base demand grows linearly, $B_i(t) = \text{base}_i \cdot (1 + g_i \cdot t/T)$, with base levels (50, 40, 60) units/day for Items 1–3.
- **Seasonality.** An annual sinusoid (amplitudes 12, 8, 15) plus an 8% weekend uplift.
- **Price paths.** Step-function list prices with roughly 60 discrete changes per product over four years, mimicking intermittent promotions.
- **Noise.** $\eta_{it} \sim \mathcal{N}(0, \sigma_i^2)$ with $\sigma = (4.0, 3.5, 5.0)$.

Table I reports the cross-price elasticity matrix used in (2). All own-price elasticities satisfy $|\varepsilon_{ii}| > 1$, so demand is elastic at the reference prices and a revenue-maximising price exists in the continuous problem. As noted in Section III, the DGP is only approximately aligned with the Poisson log-link GLM;

elasticities recovered by the GLM are therefore a sanity check against known simulator parameters, not external validation.

TABLE I: Cross-price elasticity matrix (ε_{ij}) used in the simulator. Diagonal entries are own-price elasticities; off-diagonal entries are substitution (positive) or complementary (negative) effects.

	Item 1	Item 2	Item 3
Item 1	−1.50	+0.35	−0.15
Item 2	+0.25	−1.80	+0.10
Item 3	−0.20	+0.15	−1.20

B. Hyperparameters

Table II lists the simulation and training settings used for all reported runs. The discrete action set comprises $K = 20$ prices spanning \$6.0 to \$16.0, matching the grid used in the experimental notebook. Learning parameters follow a standard DQN recipe; the econometric penalty weight is $\lambda = 30.0$ with σ_p equal to the standard deviation of the price grid. One DQN (or EG-DQN) agent is trained *per product*; there is no joint multi-product action space in this study.

TABLE II: Simulation and training hyperparameters

Parameter	Value	Notes
State dimension	5	See Section III
Action count K	20	Grid \$6.0–\$16.0
Training episodes	300	Per product
Batch size	64	Replay minibatch
Replay buffer	10,000	FIFO experience store
Discount γ	0.99	Daily pricing horizon
Learning rate α	10^{-3}	Gradient updates
Target sync	15 episodes	Hard copy $\theta^- \leftarrow \theta$
ε start/end	1.0 / 0.05	Linear decay
Masking window	episodes 1–30	EG-DQN only
λ (EG-DQN)	30.0	Reward-shaping weight
Holdout length	120 days	Item 1 revenue focus
Random seed	42	Full pipeline

C. Baseline Strategies

Five pricing strategies are compared on the same simulated trajectories and the same 120-day holdout for Item 1.

- **Fixed-Price (FP).** Charge the reference price $\bar{p}$ every day. This is a naive non-adaptive floor.
- **Rule-Based (RB).** Set price to $1.15\times$ the trailing 7-day mean of observed prices. The rule is a simple retailer heuristic, not an optimised policy.
- **GLM-Optimal (GLM).** Each day, charge $p_{\text{GLM},t}^*$ from the fitted Poisson model (Section IV). This is the pure econometric benchmark on the discrete grid.
- **Pure DQN.** The network and training procedure of Section V without state augmentation of $p_{\text{GLM},t}^*$, without reward shaping, and without action masking.
- **EG-DQN.** The full hybrid: all three integration mechanisms of Section V.

Holdout evaluation uses realised simulator demand under each policy’s chosen prices. Statistical comparisons of daily revenue series follow in the results section; the present section only fixes the protocol.

VII. EXPERIMENTS AND RESULTS

A. Demand Estimation Accuracy

We first check that the fitted Poisson GLM recovers usable demand structure in the calibrated simulator before pricing policies are compared. Figure 1 summarises the three-product series: step-function prices, seasonal demand, log-log price-demand scatters, and marginal demand distributions over the full horizon.

Table III compares estimated own-price elasticities with the ground-truth parameters of (2). Relative absolute biases are 1.2%, 1.7%, and 3.7% for Items 1–3 respectively (largest on Item 3). We therefore report elasticity recovery *within about* 3.7% in this run, rather than a tighter uniform bound. Cross-price signs match the substitution structure in Table I. Figure 2 overlays GLM and true demand curves against price: slopes agree closely with the simulator, while fitted *levels* sit above observed demand—a level misfit we treat as a modelling caveat rather than as formal structural identification.

TABLE III: GLM own-price elasticity estimates versus simulator truth (HC3-robust standard errors). Relative bias is $|\hat{\varepsilon} - \varepsilon|/|\varepsilon|$.

Product	True ε	Est. $\hat{\varepsilon}$	Bias	Rel. bias
Item 1	−1.500	−1.519	−0.019	1.2%
Item 2	−1.800	−1.770	+0.030	1.7%
Item 3	−1.200	−1.245	−0.045	3.7%




Fig. 2: GLM demand curves versus ground-truth DGP curves for Items 1–3. Annotated relative biases refer to own-price elasticity recovery (1.2%, 1.7%, 3.7%); curve *levels* are not claimed to match.

Table IV reports in-sample fit and 120-day holdout forecast accuracy for Item 1, the product used in the revenue comparison. Holdout MAPE is 5.20% and R^2 is 0.965. Figure 3 shows fitted and forecast paths against realised demand; Panel B residuals on the holdout window have mean approximately −4.57 units/day, indicating a slight systematic over-forecast even though percentage error remains modest.

TABLE IV: GLM forecast performance for Item 1 (Poisson log-link). Train covers the pre-holdout window; test is the final 120 days.

Metric	Train	Test (120-day)
MAE (units/day)	4.80	5.81
RMSE	6.17	7.07
MAPE (%)	5.22	5.20
R^2	0.983	0.965




Fig. 3: GLM demand fit and 120-day holdout forecast for Item 1 (Poisson log-link). Panel A: training actual/fitted and test actual/forecast. Panel B: test residuals (actual − forecast); mean residual ≈ -4.57 units/day.

B. Revenue Comparison

Table V and Figure 4 report holdout revenue for Item 1 over the final 120 days under the five strategies defined in Section VI. Totals are taken from Panel D of Figure 4; daily averages equal total revenue divided by 120.

On this window, **GLM-optimal and EG-DQN attain identical cumulative revenue** (\$226,764) and select the same constant price of \$6 on the discrete grid—the lower bound of the action set. The two cumulative-revenue paths in Panel A therefore overlap for the entire holdout. Relative to the fixed-price floor (\$175,651), that common outcome is +29.1% in this simulation; relative to the rule-based heuristic (\$138,864), the gap is larger still. Pure DQN earns \$183,549 (+4.5% vs. fixed-price) at an average holdout price of \$9.10—higher than the naive floor, but below the GLM/EG-DQN total. We report the Pure DQN gap descriptively here; a Mann–Whitney comparison is deferred until daily series can be re-tested cleanly.

All five learned or scripted policies are effectively constant-price on the holdout in this run (Panel C): rule-based at \$16, fixed-price at \$10, Pure DQN near \$9.10, and GLM/EG-DQN at \$6. The \$6 outcome is a *corner solution* on the discrete grid rather than rich day-to-day price variation; we return to that limitation in the Discussion. Percentage lifts versus fixed-price should therefore be read with that price path in mind. Panel B shows that daily revenue distributions for GLM-optimal and EG-DQN likewise coincide, consistent with the identical price policy.

TABLE V: Holdout revenue — Item 1, 120 days (source: Figure 4, Panel D). Daily average = total/120. “vs. FP” is relative to Fixed-Price in this simulation.

Strategy	Total (\$)	Daily (\$)	vs. FP	Avg price (\$)
Fixed-Price	175,651	1,464	—	10.00
Rule-Based	138,864	1,157	−20.9%	16.00
Pure DQN	183,549	1,530	+4.5%	9.10
GLM-Only	226,764	1,890	+29.1%	6.00
EG-DQN	226,764	1,890	+29.1%	6.00




Fig. 1: Exploratory view of the synthetic retail series (Items 1–3). Left: daily demand (solid) and list price (dashed). Centre: log–log price–demand scatter with OLS slope. Right: demand histogram with a normal overlay. The horizon spans 2021–2024 under the M5/Favorita-calibrated DGP.

C. Training Convergence

Figure 5 compares episode revenue over 300 training episodes for Pure DQN and EG-DQN on Item 1. Both agents increase cumulative episode reward under the shared ϵ -greedy schedule (right panel): first-20-episode averages are \$1,564,257 (Pure DQN) and \$1,579,526 (EG-DQN); last-20 averages rise to \$1,963,282 and \$2,003,676, improvements of +25.5% and +26.9% respectively. Best-episode rewards are likewise modestly higher for EG-DQN (\$2,017,148 vs. \$1,976,095). The moving-average trajectories in the left panel track each other closely, with EG-DQN slightly above Pure DQN for much of training. We treat this as a training-phase pattern in this run, not as evidence of faster asymptotic convergence or of holdout superiority over GLM-optimal pricing.

Logged TD losses were recorded as `nan` during training, so we do not interpret the centre panel of Figure 5 as a reliable diagnostic of learning dynamics. Guided early exploration can reduce costly exploration in nonstationary settings [13]; here that design choice sits alongside the reward-shaping and state-augmentation mechanisms rather than as a claimed variance or loss-decay result.




Fig. 5: Training trajectories for Pure DQN and EG-DQN (Item 1, 300 episodes). Left: episode revenue with a 20-episode moving average. Right: shared ϵ -greedy decay. Centre: TD-error panel is shown as exported from the notebook but is not interpreted (training logs report `Loss=nan`).

D. Ablation Study

We isolate the holdout contribution of econometric reward shaping in a three-configuration ablation on Item 1. Each variant is trained for 200 episodes under the same simulator and evaluated on the 120-day holdout; Table VI and Figure 6 summarise the outcome.

Pure DQN attains \$173,830 on the holdout window. Adding the GLM-based reward bonus (DQN+Reward) raises test revenue to \$226,764 (+30.5% relative to Pure DQN in this ablation run). The full EG-DQN configuration—state augmentation and early action masking included—matches DQN+Reward exactly on holdout revenue (\$226,764, +30.5%). In this simulation, therefore, econometric *reward*




Fig. 4: Strategy comparison on the Item 1 holdout (120 days). Panel A: cumulative revenue (GLM-Only and EG-DQN overlap). Panel B: daily revenue distributions. Panel C: holdout price paths (constant for each strategy in this run). Panel D: total revenue with printed dollar values.

shaping accounts for the principal holdout improvement over pure DQN; the additional integration channels do not increase out-of-sample revenue beyond the reward-only variant. That pattern is consistent with the view that econometric structure is most informative as a training signal rather than as a guarantee of hybrid dominance over a GLM-only benchmark [10]. Separate ablations of state augmentation alone and action masking alone are left to future work.




Fig. 6: Ablation study (Item 1). Panel A: episode revenue during training (200 episodes per configuration). Panel B: holdout test revenue; DQN+Reward and Full EG-DQN reach the same total (\$226,764).

TABLE VI: Ablation on Item 1 holdout revenue (200 training episodes per configuration). “vs. Pure DQN” is relative to the first row in this ablation run.

Variant	Test revenue (\$)	vs. Pure DQN
Pure DQN	173,830	—
DQN + Reward	226,764	+30.5%
Full EG-DQN	226,764	+30.5%

E. Statistical Significance

We test whether EG-DQN daily holdout revenue exceeds each baseline using one-sided Mann–Whitney U tests (H_1 : EG-DQN $>$ baseline) on the 120-day Item 1 window. Bootstrap 95% confidence intervals on the mean daily revenue difference use $B = 2,000$ resamples. Table VII reports three comparisons; a fourth row for Pure DQN is omitted because the exported daily series tied at the wrong revenue level in the notebook log and should be re-tested before any p -value is reported.

Relative to fixed-price and rule-based policies, EG-DQN yields significantly higher daily revenue ($p < 0.001$ in both cases). Cohen’s d values are large (2.50 and 4.63), but both naive baselines hold essentially constant prices on the holdout (Section VII), so the effect sizes reflect policy separation as much as adaptive pricing skill; we return to that point in the Discussion. Against GLM-optimal pricing, EG-DQN does *not* show higher daily revenue: $p = 0.50$, Cohen’s $d = 0.00$, and a bootstrap interval of $[+0.00, +0.00]$, consistent with the identical cumulative totals in Table V. Reporting this non-significant comparison is central to an honest reading of the hybrid design in this simulation.

TABLE VII: Statistical significance of EG-DQN daily holdout revenue vs. selected baselines (Item 1, 120 days). Mann–Whitney U , one-sided; bootstrap 95% CI on mean daily difference ($B = 2,000$).

Baseline	p -value	Cohen’s d	Bootstrap 95% CI
Fixed-Price	< 0.001	2.50	$[+421.5, +437.4]$
Rule-Based	< 0.001	4.63	$[+725.6, +751.7]$
GLM-Only	0.50	0.00	$[+0.00, +0.00]$

F. Rolling Elasticity Estimates

EG-DQN conditions learning on GLM outputs that are re-estimated as data accumulate. As a descriptive check on Item 1, we refit the Poisson log-link model on rolling twelve-month windows (advance 30 days per step) and record the own-price elasticity at each window end. Figure 7 plots the point estimate and an approximate 95% interval against the simulator parameter ($\varepsilon = -1.50$).

Over 2022–2024 the estimate moves with window composition: values lie between roughly -1.35 and -1.7 , so demand stays in the elastic range ($|\varepsilon| > 1$) throughout, but the path is not flat. The dashed truth line sits inside some intervals and outside others, which is plausible when prices change in steps and each window mixes different promotional regimes. We read this figure as evidence that the econometric input to EG-DQN is *usable* in simulation—not that elasticity is fixed, nor that the GLM prior would remain dependable in a live store without revalidation.




Fig. 7: Rolling twelve-month GLM own-price elasticity for Item 1 (30-day step). Blue line: point estimate; shaded band: approximate 95% interval; red dashed: simulator truth ($\varepsilon = -1.50$).

VIII. DISCUSSION

A. Methodological Implications

This simulation compares econometric and model-free pricing on common ground rather than treating them as

mutually exclusive toolkits. GLM-optimal pricing and EG-DQN reach the same holdout revenue on Item 1, which suggests that econometric *reward shaping* can steer a DQN toward GLM-endorsed prices when the discrete grid makes those prices optimal. The ablation points in the same direction: holdout gains over pure DQN appear once the GLM bonus is switched on, while state augmentation and masking do not add further test revenue in this run. We do not read that pattern as hybrid dominance over econometrics; it is a design observation about where GLM signals mattered under our protocol.

The distinction between model-based structure and sequential learning still matters for how one would build a pricing system. Econometric estimates supply interpretable elasticities and a price benchmark; RL supplies a policy that can, in principle, adapt when relationships drift [11]. Prior work on nonstationary RL notes that restricting early exploration can limit the cost of uninformed price trials [13]; our masking and shaping mechanisms follow that logic, though we did not establish faster convergence or lower variance from the loss logs ($\text{Loss}=\text{nan}$).

Large Cohen’s d values against fixed-price and rule-based baselines (Table VII) should be interpreted cautiously: those policies held nearly constant prices on the holdout, so daily revenue series are far apart for reasons that mix policy design with adaptive skill.

B. Limitations

Several constraints bound the findings.

Simulation scope. All evidence comes from a synthetic three-product environment with known elasticities, following the offline evaluation logic used when field tests are costly [9], [15]. We do not claim transfer to a live retailer.

Holdout tie and corner pricing. GLM-optimal and EG-DQN policies coincide at the grid minimum (\$6) on every holdout day, so the experiment does not explore interior price dynamics and the hybrid adds no out-of-sample revenue over GLM-only pricing. Percentage lifts versus fixed-price therefore partly reflect that corner outcome, not a general ranking of methods.

DGP-model alignment. The simulator is approximately consistent with a Poisson log-link GLM, which may favour econometric-guided agents relative to settings with different demand geometry. The GLM recovers own-price elasticities reasonably well yet can misstate demand *levels* (Figure 2), and holdout forecasts show a slight mean over-forecast (Figure 3).

Protocol choices. DQN training is per item, not joint across products; actions are limited to $K = 20$ prices; a single random seed and one 120-day holdout window were used. Training TD losses were not usable (nan), so learning dynamics are under-reported. The rule-based heuristic underperformed fixed-price in this run—a weakness of that heuristic, not a general statement about rule-based pricing. The reward penalty λ is not normalised to revenue units.

Incomplete ablation. State-only and mask-only variants were not estimated; rolling elasticity estimates drift over time (Figure 7), so a fixed GLM prior may need more frequent updating in practice than our simulation assumed.

C. Path to Deployment

Operational use would require steps beyond this study. A credible path would follow simulation-to-field pipelines discussed in retail RL benchmarking [9] and field-tested revenue management work [16]: calibrate a simulator or historical replay environment from store data, re-estimate demand models on a rolling basis (given time-varying elasticity in retail [6]), run fuller ablations of each EG-DQN channel, and only then consider controlled store experiments. Until those steps are completed, EG-DQN should be read as a documented hybrid *design* and evaluation protocol, not as a deployment recommendation.

IX. CONCLUSION

We described EG-DQN, a hybrid retail pricing agent that feeds GLM demand signals into a DQN through state augmentation, econometric reward shaping, and early action masking, and we evaluated the design in a synthetic three-product simulation with a 120-day holdout on Item 1. On that window, GLM-optimal pricing and EG-DQN attained the same cumulative revenue and the same price on the discrete grid; a one-sided Mann–Whitney test did not indicate higher daily revenue for EG-DQN relative to the GLM benchmark ($p = 0.50$). Relative to fixed-price and rule-based policies, EG-DQN earned more holdout revenue in this run (Mann–Whitney $p < 0.001$ in both comparisons), while pure DQN earned less than the GLM-guided policies. The ablation indicated that reward shaping carried the holdout improvement over pure DQN and that the full hybrid did not beat a reward-only variant on test revenue. Training-episode rewards were modestly higher for EG-DQN than for pure DQN.

Taken together, the study documents a reproducible benchmark and an honest comparison: econometric-guided reward shaping can align RL with GLM-optimal prices under our protocol, but we do not show holdout gains beyond GLM-only pricing. These findings are confined to a synthetic, per-item, discrete-action environment with partial alignment between the data-generating process and the GLM specification; operational deployment is not claimed.

Future work should test the design on store data, extend estimation to joint multi-product pricing [18], incorporate inventory constraints for perishable categories [19], [20], and complete ablations that isolate state augmentation and action masking. Richer price grids and field experiments will be needed before any operational recommendation can be made.

REFERENCES

- [1] P. K. Kopalle, K. Pauwels, L. Y. Akella, and M. Gangwar, “Dynamic pricing: Definition, implications for managers, and future research directions,” *Journal of Retailing*, vol. 99, no. 4, pp. 580–593, 2023. <https://doi.org/10.1016/j.jretai.2023.11.003>
- [2] R. Y. Chenavaz and S. Dimitrov, “Artificial intelligence and dynamic pricing: A systematic literature review,” *J. Appl. Econ.*, vol. 28, no. 1, p. 2466140, 2025. <https://doi.org/10.1080/15140326.2025.2466140>
- [3] J. Groeneveld *et al.*, “Reinforcement learning versus data-driven dynamic programming for finite horizon dynamic pricing markets,” *Journal of Revenue and Pricing Management*, 2025. <https://doi.org/10.1057/s41272-025-00519-8>
- [4] S. Liu *et al.*, “Data-driven dynamic pricing and inventory management of an omni-channel retailer in an uncertain demand environment,” *Expert Syst. Appl.*, vol. 244, p. 122948, 2024. <https://doi.org/10.1016/j.eswa.2023.122948>
- [5] X. Li and Z. Zheng, “Dynamic pricing with external information and inventory constraint,” *Manag. Sci.*, 2023. <https://doi.org/10.1287/mnsc.2023.4963>
- [6] W. Zhang and Y. Meng, “Dynamic modeling and forecasting of price elasticity based on time series analysis and machine learning,” *SAGE Open Economics*, 2025. <https://doi.org/10.1177/14727978251338001>
- [7] J. Bu, D. Simchi-Levi, and L. Wang, “Offline pricing and demand learning with censored data,” *Management Science*, vol. 69, no. 2, pp. 885–903, 2023. <https://doi.org/10.1287/mnsc.2022.4498>
- [8] Y. Xia, A. Arian, S. Narayanamoorthy, and J. Mabry, “RetailSynth: Synthetic data generation for retail AI systems evaluation,” *arXiv:2312.14095*, 2023.
- [9] Y. Xia *et al.*, “Simulation-based benchmarking of reinforcement learning agents for personalized retail promotions,” *arXiv:2405.10469*, 2024.
- [10] K. Safonov, “Neural network approach to demand estimation and dynamic pricing in retail,” *arXiv:2412.00920*, 2024.
- [11] M. Apte, K. Kale, P. Datar, and P. R. Deshmukh, “Dynamic retail pricing via Q-Learning — A reinforcement learning framework for enhanced revenue management,” *arXiv:2411.18261*, 2024.
- [12] P. F. Alamdar and A. Seifi, “A deep Q-learning approach to optimize ordering and dynamic pricing decisions,” *International Journal of Production Economics*, vol. 269, p. 109154, 2024. <https://doi.org/10.1016/j.ijpe.2024.109154>
- [13] W. C. Cheung and D. Simchi-Levi, “Nonstationary reinforcement learning: The blessing of (more) optimism,” *Manag. Sci.*, 2023. <https://doi.org/10.1287/mnsc.2023.4704>
- [14] Y. Zheng *et al.*, “Dual-agent deep reinforcement learning for dynamic pricing and replenishment,” *arXiv:2410.21109*, 2024.
- [15] M. Nambiar *et al.*, “Offline deep reinforcement learning for dynamic pricing of consumer credit,” *arXiv:2203.03003*, 2022.
- [16] J. Chen, Y. Xu, and P. Yu, “A reinforcement learning approach for hotel revenue management with evidence from field experiments,” *Journal of Operations Management*, vol. 69, no. 7, pp. 1176–1201, 2023. <https://doi.org/10.1002/joom.1276>
- [17] J. H. Rios and J. R. Vera, “Dynamic pricing and inventory control for multiple products in a retail chain,” *Computers & Industrial Engineering*, vol. 177, p. 109065, 2023. <https://doi.org/10.1016/j.cie.2023.109065>
- [18] K. K. N. Pillai and S. K. Amma, “Multi-agent reinforcement learning for dynamic pricing: Balancing profitability, stability and fairness,” *arXiv:2603.16888*, 2026.
- [19] T. Yavuz and O. Kaya, “Deep reinforcement learning for dynamic pricing and inventory management of perishable products,” *Applied Soft Computing*, vol. 163, p. 111864, 2024. <https://doi.org/10.1016/j.asoc.2024.111864>
- [20] Y. Nomura, Z. Liu, and T. Nishi, “Deep reinforcement learning for dynamic pricing and ordering policies in perishable inventory management,” *Appl. Sci.*, vol. 15, no. 5, p. 2421, 2025. <https://doi.org/10.3390/app15052421>